

# Sampling Estimates of Expectations

## Contents

- [Objectives](#)
- [The law of large numbers](#)

```
import numpy as np
np.set_printoptions(precision=3)
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(rc={"figure.dpi":100, 'savefig.dpi':300})
sns.set_context('notebook')
sns.set_style("ticks")
```

## Objectives

- To use the law of large numbers to estimate expectations

## The law of large numbers

### The strong law of large numbers

Take an infinite series of independent random variables  $X_1, X_2, \dots$  with the same distribution, any distribution. Such a sequence of random variables is typically called an iid sequence for *independnet identically distributed*. Let  $\mu = \mathbb{E}[X_i]$  be the common mean of these random variables. The *strong law of larger numbers* states the sampling average,

$$\bar{X}_N = \frac{X_1 + \dots X_N}{N} = \frac{1}{N} \sum_{i=1}^N X_i,$$

converges almost surely to  $\mu$  as the number of samples  $N$  goes to infinity. Mathematically, we write:

$$\bar{X}_N = \frac{1}{N} \sum_{i=1}^N X_i \rightarrow \mu \text{ a.s.}$$

The a.s. (almost surely) is a technical term from measure theory which means that the probability of this convergence happening is one.

## Demonstration with a synthetic example

Let's demonstrate the law of large numbers. We are going to take a Beta random variable:

$$X \sim \text{Beta}(\alpha, \beta),$$

where  $\alpha$  and  $\beta$  is positive numbers. We know that the expectation of the Beta is [\(see wiki\)](#):

$$\mathbb{E}[X] = \frac{\alpha}{\alpha + \beta}.$$

Let's test if the law of large numbers holds:

```
import scipy.stats as st

# Create a Beta:
alpha = 2.0
beta = 3.0
X = st.beta(alpha, beta)

# The number of samples to take
N = 5
x_samples = X.rvs(N)

# Find the real mean (expectation):
mu = X.mean()
# Find the sampling estimate of the mean:
x_bar = x_samples.mean()

# Print the results
print(f"E[X] = {mu:.4f}")
print(f"The law of large numbers with N={N:d} samples estimates it as: {x_bar:.4f}")
```

```
E[X] = 0.4000
The law of large numbers with N=5 samples estimates it as: 0.5245
```

## Questions

- Increase the number of samples  $N$  until you get closer to the correct answer.

## The Monte Carlo method for estimating integrals

Now we will use the strong law of large numbers to estimate integrals. In particular, we will start with this integral:

$$I = \mathbb{E}[g(X)] = \int g(x)p(x)dx,$$

where  $X \sim p(x)$  and  $g(x)$  is a function of  $x$ . Let  $X_1, X_2, \dots$  be independent copies of  $X$ . Then consider the random variables  $Y_1 = g(X_1), Y_2 = g(X_2), \dots$ . These random variables are also independent and identically distributed. So, the strong law of large number holds for them and we get that their sampling average converges to their mean:

$$\bar{I}_N = \frac{g(X_1) + \dots + g(X_N)}{N} = \frac{Y_1 + \dots + Y_N}{N} \rightarrow I, \text{ a.s.}$$

This is the *Monte Carlo way for estimating integrals*.

## Example: 1D expectation

Let's try it out with a test function in 1D (Example 3.4 of Robert & Casella (2004)). Assume that  $X \sim \mathcal{U}([0, 1])$  and pick:


$$g(x) = (\cos(50x) + \sin(20x))^2.$$

The correct value for the expectation can be found analytically and it is:

$$\mathbb{E}[g(x)] = 0.965.$$

```
# Define the function
g = lambda x: (np.cos(50 * x) + np.sin(20 * x)) ** 2

# Let's visualize it first
fig, ax = plt.subplots()
x = np.linspace(0, 1, 300)
ax.plot(x, g(x))
ax.set_xlabel(r"$x$")
ax.set_ylabel(r"$g(x)$");
```



Let's take the samples:

```
N = 100
x_samples = np.random.rand(N)
y_samples = g(x_samples)
```

Evaluate the sample average for all sample sizes (see [cumsum](#)).


```
I_running = np.cumsum(y_samples) / np.arange(1, N + 1)
I_running
```

```
array([1.164, 1.468, 1.107, 0.923, 0.785, 0.687, 0.717, 0.781, 0.706,
       0.635, 0.637, 0.715, 0.748, 0.844, 0.863, 0.984, 0.965, 1.008,
       0.966, 0.921, 0.878, 0.883, 0.845, 0.82 , 0.794, 0.772, 0.752,
       0.743, 0.788, 0.767, 0.746, 0.772, 0.784, 0.776, 0.791, 0.831,
       0.821, 0.802, 0.831, 0.838, 0.833, 0.814, 0.8 , 0.821, 0.862,
       0.846, 0.841, 0.83 , 0.818, 0.805, 0.791, 0.805, 0.793, 0.79 ,
       0.784, 0.77 , 0.797, 0.788, 0.778, 0.77 , 0.757, 0.747, 0.748,
       0.752, 0.765, 0.793, 0.828, 0.865, 0.891, 0.891, 0.913, 0.923,
       0.924, 0.947, 0.946, 0.944, 0.934, 0.925, 0.913, 0.903, 0.906,
       0.895, 0.912, 0.922, 0.913, 0.903, 0.894, 0.897, 0.888, 0.913,
       0.928, 0.918, 0.91 , 0.907, 0.898, 0.889, 0.889, 0.883, 0.875,
       0.901])
```

Plot this:

```
fig, ax = plt.subplots()
ax.plot(np.arange(1, N+1), I_running)
ax.plot(np.arange(1, N+1), [0.965] * N, color="r")
ax.set_xlabel(r"$N$")
ax.set_ylabel(r"$\bar{I}_N$");
```

syntax for plotting at 0.965 for each discrete x-location?



Objective is to achieve convergence as cumulative sum consists of more samples from the rv

## Questions

- Increase  $N$  until you get an answer that is close enough to the correct answer (the red line).
- Reduce  $N$  back to a small number, say 1,000. Run the code 2-3 times to observe that every time you get a slightly different answer...

---

By Ilias Billionis (ibillion[at]purdue.edu)

© Copyright 2021.