

# Linear regression with a single variable

## Contents

- [Objectives](#)
- [An example where things work as expected](#)
- [Questions](#)
- [An example where things do not work as expected: underfitting](#)
- [Questions](#)

```
import numpy as np
np.set_printoptions(precision=3)
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(rc={"figure.dpi":100, "savefig.dpi":300})
sns.set_context("notebook")
sns.set_style("ticks")
```

## Objectives

- To introduce linear regression with a single variable

## An example where things work as expected

Let's create a synthetic dataset to introduce the basic concepts. It has to be synthetic because we want to know what the ground truth is. Let's start with pairs of  $x$  and  $y$  which definitely have a linear relationship, albeit  $y$  may be contaminated with Gaussian noise. In particular, we generate the data from:

$$y_i = -0.5 + 2x_i + 0.1\epsilon_i,$$


where  $\epsilon_i \sim N(0, 1)$  and where we sample  $x_i \sim U([0, 1])$ . Here is how to generate this synthetic dataset and how it looks like.

```
np.random.seed(12345)

num_obs = 10
x = np.random.rand(num_obs)
w0_true = -0.5
w1_true = 2.0
sigma_true = 0.1
y = (
    w0_true
    + w1_true * x
    + sigma_true * np.random.randn(num_obs)
)
```

Let's plot the data:

```
fig, ax = plt.subplots()
ax.plot(x, y, 'x', label='Observed data')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')
plt.legend(loc='best');
```



We will now use least squares to fit the data to this linear model:

$$y = w_0 + w_1 x.$$

As we discussed in the previous section, least squares minimize the square loss:

$$L(\mathbf{w}) = \sum_{i=1}^N (y_i - w_0 - w_1 x_i)^2 = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2,$$

where  $\mathbf{y} = (y_1, \dots, y_N)$  is the vector of observations,  $\mathbf{w} = (w_0, w_1)$  is the weight vector, and the  $N \times 2$  **design matrix**  $\mathbf{X}$  is:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}.$$

We definitely need to make the design matrix  $\mathbf{X}$ :

```
# Put together a column of ones next to the observed x's
X = np.hstack(
    [np.ones((num_obs, 1)), x.reshape((num_obs, 1))]
)
X
```

```
array([[1.    , 0.93  ],
       [1.    , 0.316],
       [1.    , 0.184],
       [1.    , 0.205],
       [1.    , 0.568],
       [1.    , 0.596],
       [1.    , 0.965],
       [1.    , 0.653],
       [1.    , 0.749],
       [1.    , 0.654]])
```

Once we have this, we can use [numpy.linalg.lstsq](https://numpy.org/doc/stable/reference/generated/numpy.linalg.lstsq.html) to solve the least squares problem. This function solves in a smart way the linear system we derived in the previous section, i.e.,

See this page for a lot more information on linear regression:  
[https://en.wikipedia.org/wiki/Least\\_squares](https://en.wikipedia.org/wiki/Least_squares)

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}.$$

Start point:  $y = Xw$ , isolate for  $w$

It works as follows:

```
w, _, _, _ = np.linalg.lstsq(X, y, rcond=None)
print(f'w_0 = {w[0]:1.2f}')
print(f'w_1 = {w[1]:1.2f}')
```

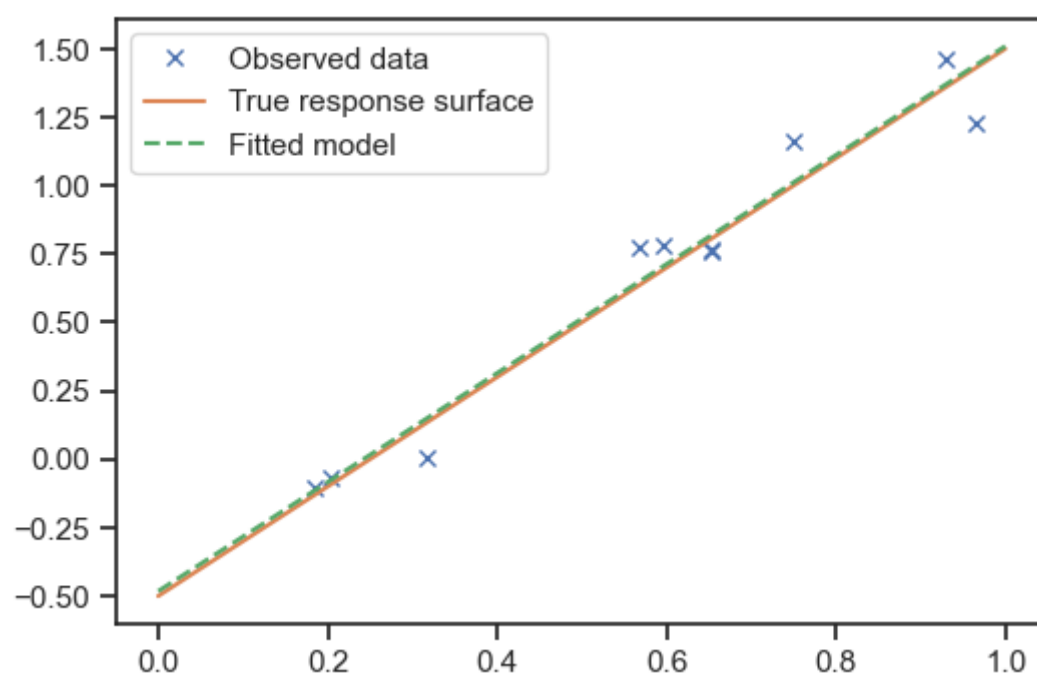
```
w_0 = -0.48
w_1 = 1.99
```

So, you see that the values we found for  $w_0$  and  $w_1$  are close to the correct values, but not exactly the same. That is fine. There is noise in the data and we have only used ten observations. **The more noise there is, the more observations it would take to identify the regression coefficients correctly.**

Let's now plot the regression function against the data:

```
# Make predictions
# Some points on which to evaluate the regression function
xx = np.linspace(0, 1, 100)
# The true connection between x and y
yy_true = w0_true + w1_true * xx
# The model we just fitted
yy = w[0] + w[1] * xx

# Plot them
fig, ax = plt.subplots()
# plot the data again
ax.plot(x, y, 'x', label='Observed data')
# overlay the true
ax.plot(xx, yy_true, label='True response surface')
# overlay our prediction
ax.plot(xx, yy, '--', label='Fitted model')
plt.legend(loc='best');
```



## Questions

- Try increasing `num_obs` to 100. Does the fit improve? Conclusion: When you training with least squares, the more data you have the better.
- Try decreasing `num_obs` to 2. What is happening here? This is an example of fitting the noise.

## An example where things do not work as expected: underfitting

Let's try to fit a linear regression model to data generated from:

$$y_i = -0.5 + 2x_i + 2x_i^2 + \epsilon_i,$$

where  $\epsilon_i \sim N(0, 1)$  and where we sample  $x_i \sim U([-1, 1])$ :


```

np.random.seed(12345)

num_obs = 10
x = -1.0 + 2 * np.random.rand(num_obs)  # method to generate uniform random numbers between -1 and 1
w0_true = -0.5
w1_true = 2.0
w2_true = 2.0
sigma_true = 0.1
y = (
    w0_true
    + w1_true * x
    + w2_true * x ** 2
    + sigma_true * np.random.randn(num_obs)
)

fig, ax = plt.subplots()
ax.plot(x, y, 'x', label='Observed data')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')
plt.legend(loc='best');

```



We will still try to fit a linear model to this dataset. We know that it is not going to work well, but let's try it anyway. First, create the design matrix just like before:

```

X = np.hstack(
    [np.ones((num_obs, 1)), x.reshape((num_obs, 1))]
)

w, _, _, _ = np.linalg.lstsq(X, y, rcond=None)

print(f'w_0 = {w[0]:1.2f}')
print(f'w_1 = {w[1]:1.2f}')

```

```

w_0 = 0.03
w_1 = 2.46


```

```

# Make predictions
xx = np.linspace(-1, 1, 100)
yy_true = w0_true + w1_true * xx + w2_true * xx ** 2  # degree 2
yy = w[0] + w[1] * xx  # degree 1

# Plot them
fig, ax = plt.subplots()
ax.plot(x, y, 'x', label='Observed data')
ax.plot(xx, yy_true, label='True response surface')
ax.plot(xx, yy, '--', label='Fitted model')
plt.legend(loc='best');

```



## Questions

- Experiment with very small `num_obs`. If you did not know what the true response surface was, would you be able to say whether or not the fit is good?
- Experiment with a big `num_obs`. Does the fit improve? This is an example of *underfitting*. Your model does not have enough expressivity to capture the data.

---

By Ilias Bilionis (ibilion[at]purdue.edu)

© Copyright 2021.